

CMSC320 – Introduction to Data Science

Final Tutorial

Spring Semester 202k

José Manuel Calderón Trilla and Elias Gonzalez
University of Maryland
Department of Computer Science

Motivation

There will be no final exam for CMSC320, instead students are asked to submit a *tutorial* that walks the reader through the Data Science pipeline. The subject matter of this tutorial is far less important than the ability to communicate the approach throughout and a meaningful discussion of the implications/interpretations of the final results.

For the purposes of this tutorial, we will assume that ‘The Data Science Pipeline’ has the following phases:

1. Data collection/curation + parsing (if necessary)
2. Data management/representation
3. Exploratory data analysis
4. Hypothesis testing
5. Communication of insights attained

Each of these stages must be present in some form in your submission.

It is required that each tutorial is a *self-contained* artifact, using a combination of Markdown and Python code within a Jupyter Notebook. This artifact should be publicly available on the web. Github Pages is a reasonable choice for this, but it is not required (Google Collab is not considered *static* but I believe it is possible to extract a static page from that system). We will discuss this later in this document.

The dataset chosen should be publicly available (so that we can replicate your results). Some possible sources of data:

1. A curated list of datasets: <https://github.com/awesomedata/awesome-public-datasets>
2. The U.S. Government is a fantastic source of open data: <https://www.data.gov/>
3. Often individual States will also host their own open data, here is the link for Maryland: <https://opendata.maryland.gov/>
4. Microsoft has a list of open datasets: <https://azure.microsoft.com/en-us/services/open-datasets/catalog/>
5. The National Institute of Health provides many datasets, here are the COVID related datasets: <https://datascience.nih.gov/covid-19-open-access-resources>

1 Expectations

In general we would expect a good submission to provide the following, at a minimum:

- 1500+ words of prose in English, describing the process throughout and a discussion of the insights attained
- Approximately 150 lines of non-contrived Python
- Well-labelled figures showing important aspects of the analysis
- Links to external documentation and resources that would be useful in understanding the approach.

1.1 Groups

Groups are allowed for the final project, up to 3 people per group. Members of a group do not have to be enrolled in the same section of CMSC320. As the group size increases, so would the scale of the expectations: more people should result in a more thorough tutorial. Groups must be pre-registered by April 16th (one month before the final deadline). In order to prevent miscommunication about whether students are part of a group or not, all group members must register independently. We will provide a form for students to register their groups and announce the form to the class.

The graders recognize that there are some aspects of the data science pipeline where there may only be one ‘right’ way to do something. For example, if you’re accessing data from a pre-populated SQL database, we do not expect a group to discuss multiple ways of accessing the data! However, there are many aspects in data-visualization, data-exploration, or data-analysis (just as examples), where there are multiple methods of accomplishing the same task, we would expect a group to discuss more of those alternatives. Often this results in only small amount of extra code, but a significant amount of extra discussion.

2 Examples

The following are links to final projects from past semesters. They should be seen as a rough guide to what is expected and to the variety of topics that can be pursued and not as examples of the highest-scoring submissions.

- The golden age of rap: <http://rstumbaugh.me/hiphop-analysis/>
- Predicting a win in Rainbow Six: Siege: <https://jiglesia3.github.io/>
- What makes the best defensive footballers? <https://bdaisey.github.io/>
- Maryland and peer institutions’ faculty/student counts: <https://krixly.github.io/>.
- Analysis of crime data in College Park: <https://andresgogo.github.io/>

3 The Deliverable

We recommend ‘GitHub Pages’ (<https://pages.github.com/>) or TerpConnect (<https://terpconnect.umd.edu/we>) for hosting your final tutorial. These services are both provided, free of charge, by their respective organizations. Neither of these services try to be ‘smart’ about hosting a Jupyter Notebook, it only hosts a static HTML page. So generating that HTML can be done independently. This also has the side benefit that there is no risk of data sources shifting or changing, causing your project to fail, all data manipulation, analyses, and figures are computed ahead of time and hosted statically.

In order to use GitHub Pages, the following 3 items are required:

1. A GitHub account
2. A repository named `<account-name>.github.io`, where `<account-name>` is your unique GitHub account handle
3. An HTML export of your final tutorial

You can use the same repository for storing and versioning your final project, in fact I recommend it!), but that is not a require of GitHub pages itself. There are many tutorials online for using GitHub pages. Personally, I have found the official page to be sufficient: <https://pages.github.com/>.

3.1 Format of your deliverable

The formatting for the majority of the deliverable is left to your discretion. However, each submission must begin with the title of the tutorial, providing a rough idea of the topic, followed by your name (and all members of the group).

3.2 Submission of your tutorial

There are two deadlines for the final tutorial:

- **4PM EDT May 9th, 2022:** Your site being accessible via the web (but your tutorial not yet completed)
- **4PM EDT May 16th, 2022:** Your tutorial being completed, accesible via the web, and able to be scrapped

The submission to staff is simply the URL for your hosted tutorial. All tutorials must be completed by the University’s specified Final Exam time: **4:00PM EDT on May 16th 2022**. Due to the short time-table for final grades, no late submissions will be accepted. At the time of the deadline, a mirroring script will be run, downloading all of the tutorials. As such, no changes to the site after the deadline will be considered as part of the submission.

4 Assessment

The following dimensions of each submission will be given a rating between 1-10:

1. Motivation
2. Understanding
3. Resources
4. Prose
5. Code
6. Communication of Approach
7. Formatting and Subjective Evaluation

In general, the tutorial should contain at least 1500 words of prose and 150 lines of (non-padded, legitimate) Python code, along with appropriate documentation, visualization, and links to any external information that might help the reader. You are welcome to do this project individually or in a group of size at most three; we'll scale up the expectations accordingly as group size increases.

4.1 Github Pages

GitHub provides a service called Pages (<https://pages.github.com/>) that provides website hosting functionality backed by a GitHub-based git repository. We would like you to host your final project on a GitHub Pages project site. To do this, you will need to:

1. Create a GitHub account (or use the one you already have) with username **username**.
2. Create a git repository titled **username.github.io**; make sure **username** is the same as whatever you chose for your global GitHub account.
3. Create a project within this repository. This is where you'll dump your iPython Notebook file and an HTML export of that Notebook file.

These instructions are also given directly on the front page of <https://pages.github.com/>; following those instructions should be fine!

Motivation: each tutorial should be sufficiently motivated. If there is not motivation for the analysis, why would we 'do data science' on this topic?

Understanding: the reader of the tutorial should walk away with some new understanding of the topic at hand. If it's not possible for a reader to state 'what they learned' from reading your tutorial, then why do the analysis?

Resources: tutorials should help the reader learn a skill, but they should also provide a launching pad for the reader to further develop that skill. The tutorial should link to additional resources wherever appropriate, so that a well-motivated reader can read further on techniques that have been used in the tutorial.

Prose: it's very easy to write the literal English for what the Python code is doing, but that's not very useful. The prose should enhance, the tutorial, adding additional context and insight.

Code: code should be clear and commented. Function definitions should be described and given context/motivation. If the prose helps the reader understand *why* you've written the code, the comments in the code should be sufficient for the reader to learn how.

Pipeline: all stages of the pipeline should be discussed. We will be looking for 'good science', with discussion of each stage and what its implications/consequences are.

Communication of Approach: every technical choice has alternatives, why did you choose the approach taken in the tutorial? A reader should walk away with some idea of what the trade-offs may be.

Formatting and Subjective Evaluation: does the tutorial seem polished and 'publishable', or haphazard and quickly thrown together? The tutorials should read as well put-together and having undergone a few iterations of editing and refinement. This should be the easiest of the dimensions.

Dimension	Points Available
Motivation	10
Understanding	10
Further Resources	10
Prose	20
Code	10
Pipeline	10
Communication of Approach	20
Subjective Evaluation	10
Total	100